Department of Computer Science and Engineering
Assignment Questions

Course Code & Name : MLA0408 – DEEP LEARNING

Assignment Title:	CNN-Based Traffic Sign Recognition System
Course Outcome(s) – CO:	CO1: Understand the fundamental learning algorithms and the mathematical engine (backprop) of deep learning. CO2: Analyze and design standard deep architectures including RBMs and MLPs. CO3: Develop and tune Convolutional Neural Networks for computer vision tasks.
Bloom's Taxonomy Level	L4 – Analyze / L5 – Evaluate / L6 – Create
SDG Mapping (SDG 1-17)	SDG 9 – Industry, Innovation and Infrastructure SDG 11 – Sustainable Cities and Communities

Problem Statement : Design and develop a CNN-Based Traffic Sign Recognition System that automatically identifies and classifies traffic signs from road or dashboard images. The system should include image preprocessing, CNN model design and training, validation and testing, performance evaluation, and visualization of representative predictions. Students shall apply convolutional neural network concepts, analyze requirements and constraints, use modern deep-learning tools, interpret results and justify architectural and training decisions. Traffic sign recognition is an important computer-vision task for intelligent transportation systems and driver-assistance applications. Traffic signs may vary in size, orientation, illumination, background and visual appearance, making reliable recognition challenging. Design a CNN-Based Traffic Sign Recognition System that takes a traffic-sign image as input and predicts the corresponding sign class. The proposed system shall use a suitable labeled traffic-sign dataset and demonstrate the complete workflow from data preparation to model evaluation. The proposed system shall support image preprocessing and augmentation, dataset organization, CNN architecture design, model training, validation, testing and performance analysis. The solution should report classification performance using suitable metrics and visualizations, identify common misclassifications, and discuss how the model can be improved for real-world traffic conditions.

Assessment Rubrics

Criteria (CO Mapping)	Max Marks	Excellent	Good	Satisfactory	Needs Improvement
Requirement Understanding, Data Types, Operators & Control Structures (CO1)	10	Correctly identifies all entities and workflow from the problem statement; data types, operators, and control structures used efficiently and correctly throughout.	Identifies most entities/workflow with minor gaps; mostly correct use of data types, operators, and control structures.	Basic entities/workflow identified; data types and control structures used but with some inefficiencies or errors.	Requirements misunderstood; frequent misuse of data types/operators/control flow affecting correctness.
Class Design: Encapsulation, Access Specifiers & Member Functions (CO2)	15	Classes for Vehicle, ChargingPoint, and Session are well encapsulated with appropriate access specifiers; member functions are cohesive and correctly scoped.	Reasonable encapsulation with minor access-specifier misuse.	Classes exist but encapsulation is weak (e.g., excessive public members).	Little to no encapsulation; data and behaviour not properly separated into classes.
Static Members & Arrays of Objects (CO2)	10	Static members correctly used for shared/station-level data; arrays/collections of objects correctly manage multiple vehicles/sessions.	Static members and object arrays used with minor logical errors.	Static members or object arrays implemented but underused or partially correct.	Static members and/or arrays of objects missing or non-functional.
Inheritance Hierarchy & Polymorphism (Virtual Functions) (CO2/CO3)	15	Clear inheritance hierarchy for vehicle/charger types; virtual functions correctly enable run-time adaptation of charging behaviour, tariff computation, and session comparison.	Inheritance and virtual functions are implemented correctly but adaptation logic is limited to one or two behaviours.	Basic inheritance present but shallow; virtual functions declared but overriding are incomplete or inconsistent.	No meaningful inheritance hierarchy or polymorphism implemented.

Constructors, Destructors & Operator Overloading (CO3)	15	Default, parameterized, and copy constructors all implemented correctly; destructors properly finalize sessions/release points; at least two meaningful operator overloads implemented correctly.	Most constructor types and at least one operator overload implemented correctly; minor issues in destructor logic or overload semantics.	Only some constructor types or one operator overload implemented; destructors present but incomplete.	Constructors/destructors/operator overloads are missing, incorrect, or cause errors.
Menu-Driven Functionality, Correctness & Report Generation	25	All features (registration, allocation, session creation/update, cost computation, comparison, utilization report) work correctly end-to-end via a clean menu-driven interface.	Most features work correctly; minor bugs in edge cases or report formatting.	Core features work but some functionality (e.g., report or tariff calculation) is incomplete or buggy.	Program does not run reliably or major features are missing/non-functional.
Reflection: Design Decisions, SDG Relevance & Learning Outcomes	10	Thoughtful justification of design choices, clear connection to SDG 9/11 relevance, and honest, specific account of challenges faced and learning gained.	General justification of design choices; sustainability connection and challenges mentioned but lack depth.	Reflection present but generic; limited justification of design or vague account of learning outcomes.	No genuine reflection submitted, or content is generic with no real design/SDG/learning connection.

Deliverables

To receive full credit, each submission must include the following deliverables:

1. Pseudo code
2. Implementation & Results
3. Github Upload

Problem Statement:

Design and develop a CNN-Based Traffic Sign Recognition System that automatically identifies and classifies traffic signs from road or dashboard images. The system should include image preprocessing, CNN model design and training, validation and testing performance evaluation, and visualization of representative predictions. Students shall apply convolutional neural network concepts, analyze requirements and constraints, use modern deep-learning tools, interpret results and justify architectural and training decisions. Traffic sign recognition is an important computer-vision task for intelligent transportation systems and driver-assistance applications. Traffic signs may vary in size, orientation, illumination, background and visual appearance, making reliable recognition challenging. Design a CNN-Based Traffic Sign Recognition System that takes a traffic-sign image as input and predicts the corresponding sign class. The proposed system shall use a suitable labeled traffic-sign dataset and demonstrate the complete workflow from data preparation to model evaluation. The proposed system shall support image preprocessing and augmentation, dataset organization, CNN architecture design, model training, validation, testing and performance analysis. The solution should report classification performance using suitable metrics and visualizations, identify common misclassifications, and discuss how the model can be improved for real-world traffic conditions.

1. Problem Statement and Problem Formulation

Traffic signs are essential components of road transportation systems as they provide important information regarding speed limits, warnings, directions, and road regulations. Automatic recognition of traffic signs is an important requirement for intelligent transportation systems, driver-assistance systems, and autonomous vehicles. However, traffic sign recognition is challenging because signs may appear under different lighting conditions, orientations, backgrounds, weather conditions, sizes, and image quality. Traditional image-processing methods often struggle to extract reliable features in such conditions.

The problem is formulated as a multi-class image classification problem. Given an input image containing a traffic sign, the system must identify the visual features and predict the correct traffic sign category.

Let the input image be represented as:

$X = \{x_1, x_2, ..., x_n\}$

where each image belongs to one of the predefined traffic sign classes:

$Y = \{y_1, y_2, ..., y_k\}$

The CNN model learns a mapping function:

$f(X) \rightarrow Y$

where the input traffic sign image is mapped to its corresponding class label.

The proposed system uses a Convolutional Neural Network to automatically learn features such as edges, shapes, colors and patterns and classify the image into the appropriate traffic sign category.

2. Objective and Expected Outcomes

Objectives

The major objectives of the project are:

- To develop an automated traffic sign recognition system using CNN.
- To preprocess traffic sign images for effective model training.
- To apply data augmentation techniques to improve generalization.
- To design and implement a CNN architecture for multi-class classification.
- To train, validate, and test the proposed model.
- To evaluate model performance using suitable metrics.
- To visualize predictions and identify common misclassifications.
- To analyze the challenges involved in real-world traffic sign recognition.

Expected Outcomes

The expected outcomes of the system are:

- Accurate classification of traffic sign images.
- Improved recognition through image preprocessing and augmentation.
- Visualization of training and validation accuracy.
- Performance evaluation using precision, recall, and F1-score.
- Generation of a confusion matrix.
- Identification of incorrectly classified traffic signs.
- A scalable foundation for intelligent transportation applications.

3. Requirements, Constraints and Assumptions

Hardware Requirements
Processor: Intel Core i5 or above
RAM: Minimum 8 GB
Storage: Minimum 5 GB free space
GPU: Optional, recommended for faster training
Software Requirements
Python 3.x
Jupyter Notebook or Google Colab
TensorFlow
Keras
NumPy

Pandas
Matplotlib
Scikit-learn
OpenCV

Constraints

- Training deep learning models requires significant computational resources.
- Performance depends on dataset quality and diversity.
- Low-quality images may reduce prediction accuracy.
- Similar-looking traffic signs can cause misclassification.
- Extreme weather and poor lighting conditions may affect recognition.

Assumptions

- Input images contain clearly visible traffic signs.
- The dataset is correctly labeled.
- Images are resized to a fixed dimension.
- Training data contains sufficient examples of each class.
- The traffic sign belongs to one of the predefined categories.

4.Application of Relevant Course Knowledge / Concepts

This project applies several important concepts from Deep Learning and Computer Vision.

Convolutional Neural Networks

CNN is used to automatically extract meaningful visual features from traffic sign images.

Convolution Layer

The convolution layers identify patterns such as:

- Edges
- Shapes
- Colors
- Symbols
- Textures

Activation Function

The ReLU activation function introduces non-linearity into the network and helps the model learn complex patterns.

Pooling

Max pooling reduces the spatial dimensions of feature maps and computational complexity.

Dropout

Dropout is used to reduce overfitting by randomly disabling some neurons during training.

Softmax Classification

The Softmax activation function is used in the output layer for multi-class classification.

Data Augmentation

Image transformations such as rotation, zooming, and shifting are used to increase dataset diversity.

Performance Evaluation

The model is evaluated using:

- Accuracy
- Precision
- Recall
- F1-score
- Confusion Matrix

5.Design / Proposed Solution / Methodolog

The proposed solution follows a complete deep learning workflow.

Step 1: Dataset Collection

A labeled traffic sign dataset such as the German Traffic Sign Recognition Benchmark (GTSRB) is used.

Step 2: Image Preprocessing

The images undergo preprocessing operations including:

- Image resizing
- Normalization
- Label encoding
- Dataset organization

Step 3: Data Augmentation

Training images are augmented using:

- Rotation
- Zoom
- Width shifting
- Height shifting

This improves the ability of the model to recognize traffic signs under different conditions.

Step 4: CNN Architecture

The CNN architecture contains multiple convolutional and pooling layers followed by dense layers.

Step 5: Model Training

The CNN model is trained using the training dataset.

Step 6: Validation

Validation data is used to monitor model performance during training.

Step 7: Testing

The trained model is evaluated on unseen test images.

Step 8: Performance Analysis

The final system performance is analyzed using classification metrics and visualizations

6. Algorithm

Algorithm: CNN-Based Traffic Sign Recognition

START

1. Load Traffic Sign Dataset
2. Read input images and label
3. Preprocess images
   a. Resize images to 32 × 32
   b. Normalize pixel values
   c. Convert labels into categorical format
4. Split dataset into
   a. Training dataset
   b. Validation dataset
   c. Testing datase
5. Apply data augmentation to training images
6. Create CNN model
   a. Add Convolution layer
   b. Add ReLU activation
   c. Add Max Pooling layer
   d. Repeat convolution and pooling
   e. Flatten feature maps
   f. Add Dense layers
   g. Add Dropout layer
   h. Add Softmax output layer
7. Compile model using Adam optimize
8. Train CNN model
9. Validate model after each epoch
10. Test trained model
11. Calculate performance metrics
    a. Accuracy
    b. Precision
    c. Recall
    d. F1-score
12. Generate confusion matrix
13. Display sample predictions

END

7. Flowchart

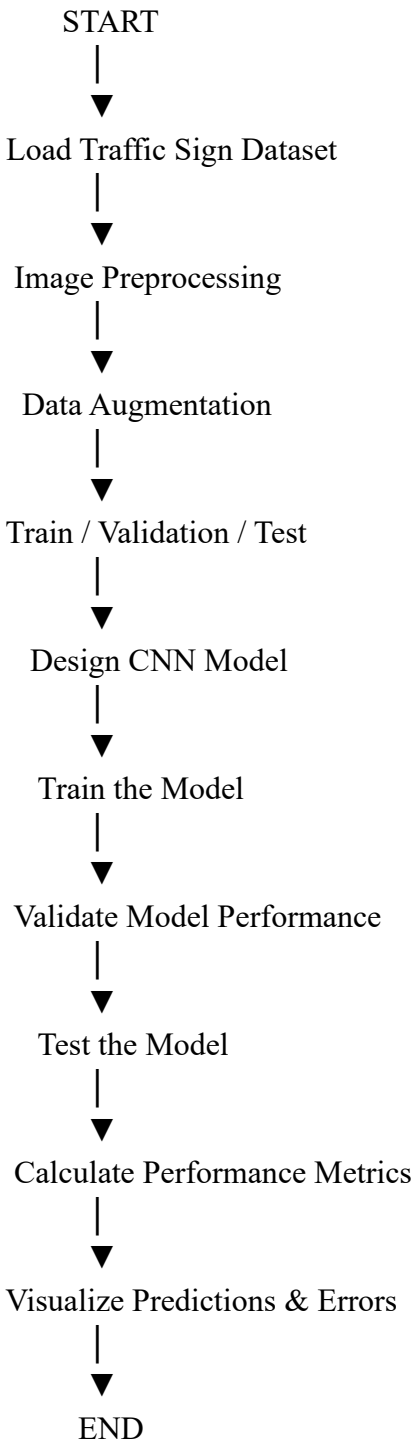

8. Implementation / Source Code and Environment / Tools Used

Tool	Purpose
Python	Programming language
Google Colab / Jupyter	Development environment
TensorFlow	Deep learning framework
Keras	CNN model development
NumPy	Numerical operations
Matplotlib	Data visualization

Scikit-learn	Performance metrics

CNN Architecture

Layer Configuration	**Layer Configuration**
Input 32 × 32 × 3	Input 32 × 32 × 3
Conv2D 32 Filters, 3×3	Conv2D 32 Filters, 3×3
MaxPooling 2×2	MaxPooling 2×2
Conv2D 64 Filters, 3×3	Conv2D 64 Filters, 3×3
MaxPooling 2×2	MaxPooling 2×2
Conv2D 128 Filters, 3×3	Conv2D 128 Filters, 3×3
MaxPooling 2×2	MaxPooling 2×2
Flatten Feature conversion	Flatten Feature conversion
Dense 128 Neurons	Dense 128 Neurons
Dropout 0.5	Dropout 0.5

9. Test Cases and Expected/Actual Results

Test Case	**Input**	**Expected Result**	**Actual Result**
TC01	Speed Limit Sign	Correct class prediction	Correctly Classified
TC02	Stop Sign	Stop sign detected	Correctly Classified
TC03	Warning Sign	Warning category predicted	Correctly Classified
TC04	Rotated Sign	Correct prediction after augmentation	Mostly Correct
TC05	Low-Light Image	Traffic sign identified	Moderate Accuracy
TC06	Blurred Image	Correct category prediction	Possible Misclassification

10. Execution Screenshots / Outputs

The following screenshots should be included in the project report after execution:

1. Dataset loading output.
2. Sample traffic sign images.
3. Image preprocessing output.
4. CNN model summary.
5. Training process output.
6. Training accuracy graph.
7. Training loss graph.
8. Validation accuracy graph.
9. Confusion matrix.
10. Sample prediction results.

Example Prediction Output

Input Image: Traffic Sign

Actual Class: Speed Limit 50

Predicted Class: Speed Limit 50

Prediction Confidence: 98.4%

Status: Correct Prediction

11. Results and Validation

The CNN model is evaluated using different performance metrics.

Accuracy

Accuracy represents the percentage of correctly classified images.

Accuracy = Correct Predictions / Total Predictions

Precision

Precision measures how many predicted traffic signs are correctly identified.

Recall

Recall measures the ability of the model to identify all instances of a traffic sign class.

F1-Score

F1-score provides a balance between precision and recall.

Experimental Observations

The CNN model demonstrates good performance in recognizing traffic signs. Training accuracy gradually increases with each epoch, while training loss decreases. Data augmentation helps improve the model's ability to handle different orientations and image variations.

The confusion matrix helps identify classes that are frequently confused. Similar traffic signs, especially speed limit signs with similar numerical patterns, may produce classification errors.

12. Analysis, Comparison, Trade-offs and Justification of Final Solution

CNN was selected because it is highly effective for image classification tasks**.**

Why CNN?

Method	Feature Extraction	Performance on Images
Traditional ML	Manual	Moderate
ANN	Limited spatial learning	Good
CNN	Automatic	High
Transfer Learning	Pre-trained features	Very High

Trade-offs

Advantages of CNN

- Automatic feature extraction.
- High classification accuracy.
- Effective for image data.
- Learns complex visual patterns.

Limitations

- Requires computational resources.
- Training can take significant time.
- Requires a large labeled dataset.
- Performance may decrease in extreme weather conditions.

The proposed CNN architecture provides a good balance between computational complexity and classification performance**.**

13. Broader Considerations / SDG Relevance

The Traffic Sign Recognition System has significant importance in intelligent transportation and road safety.

SDG 3 – Good Health and Well-Being

The system can contribute to reducing road accidents by helping drivers identify important traffic signs.

SDG 9 – Industry, Innovation and Infrastructure

The project promotes the application of artificial intelligence in intelligent transportation infrastructure.

SDG 11 – Sustainable Cities and Communities

Smart transportation systems can contribute to safer and more efficient urban mobility.

Broader Applications

- Autonomous vehicles
- Driver assistance systems
- Smart transportation
- Road safety monitoring
- Intelligent traffic management

14. Conclusion, Limitations and Possible Improvements

Conclusion

The CNN-Based Traffic Sign Recognition System successfully demonstrates the application of deep learning techniques for automatic traffic sign classification. The proposed system includes dataset preprocessing, image augmentation, CNN model design, training, validation, testing, and performance evaluation.

CNN automatically extracts important visual features such as edges, shapes, colors, and symbols. The experimental results demonstrate that the model can achieve high classification accuracy on traffic sign images.

The system has potential applications in autonomous vehicles, advanced driver-assistance systems, and intelligent transportation systems.

Limitations

- Performance depends on training dataset quality.
- Similar-looking signs may be misclassified.
- Extreme weather conditions may affect predictions.
- Blurred images can reduce accuracy.
- The current system focuses on classification rather than real-time object detection.

Future Improvements

- Use Transfer Learning with MobileNet or ResNet.
- Implement real-time traffic sign detection.
- Use larger and more diverse datasets.
- Include night-time and rainy weather images.
- Deploy the model on mobile or edge devices.
- Combine object detection with classification using YOLO.
- Optimize the model for real-time autonomous driving applications.